

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy and how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargated, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).



What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).



Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies

- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending

a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and twitter.com.

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to twitter.com on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on twitter.com, iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: 0.

In the top menu, tap your **profile** icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

0. In the top menu, you will either see a **navigation menu** icon or your **profile**

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

[Username squatting policy](#)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

[Child sexual exploitation policy](#)

[Twitter removes content promoting child sexual exploitation](#)

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at [/forms/platform](#).

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spammy or duplicative use of mentions and replies may result in enforcement action, such as the removal of

your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your

privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.
