

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;

- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and

- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local

law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name,

or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide

Twitter Rules: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;

- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing

their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
 - text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
 - information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;

- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order

to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered,

slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity’s name, please review Twitter’s [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

**Media not displayed**

This image has been removed in response to a report from the copyright holder.

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the

material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;

- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks

to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.